4.5 Years as Data Engineer at KPIT

Rahul Waghmare

7841098622

rahulsadu95@gmail.com

India, Pune

Professional summary

4.5 Years as Data Engineer at KPIT from july 2021 to till date

Data Engineer with 4.5+ years of experience in Data Engineering, Big Data technologies and Data Analysis, . Proficient in Python, PySpark, Databricks SQL, Hadoop, Hive , Azure Cloud with expertise in building and managing scalable data pipelines using Azure Data Factory. Skilled in Data Mining, Data Preparation, Data modeling and ETL processes, ensuring high-quality data flow for analytical and business needs. Experienced in handling large datasets, implementing Machine Learning algorithms, and working with cloud platforms for data storage and processing. Strong knowledge of Proof of Concepts (PoC) and gap analysis to drive data-driven solutions for enterprise applications.

Experience

Data Engineer

Maersk /

Project Name

Cargo Capacity Planning Platform

Technical Skills

Python (pandas, openpyxl, xlrd, deep-translator, logging), PySpark (Spark SQL, functions, data types), Delta Lake, Azure Logic Apps, Azure Data Factory, ADLS Gen2, Azure Blob Storage, Azure Databricks, Unity Catalog, Azure ML Studio, Couchbase, Azure Key Vault, CI/CD & DevOps scripting

Project Overview

This project focused on building a scalable, cloud-native data platform to optimize cargo utilization for Walwil and its automotive partners in the shipping and logistics domain. The organization faced fragmented data sources, inconsistent schemas, multilingual inputs, and manual processing of files received via email and SharePoint, which delayed analytics and blocked ML-driven optimization.

The solution automated end-to-end ingestion, validation, transformation, and delivery of customer shipment data using Microsoft Azure and Databricks. A configuration-driven approach enabled rapid onboarding of 40+ clients, while Medallion Architecture (Bronze–Silver–Gold) ensured structured, high-quality data processing. Curated Gold-layer datasets were published to Couchbase, making the data ML-ready for cargo optimization models and advanced analytics.

Roles

- Designed and implemented automated ingestion pipelines using Azure Logic Apps for email-based files and SharePoint API integrations
- Built Azure Data Factory pipelines for schema validation, normalization, multilingual data standardization, and exception handling
- Developed configuration-driven pipelines enabling dynamic onboarding and processing for 40+ automotive clients
- Implemented Medallion Architecture in Databricks using PySpark and Delta Lake (Bronze, Silver, Gold layers)
- Enforced data quality checks, deduplication, quarantine handling, and audit-ready logging
-

Education

bachelor of engineering

2007 - 2011

sres coe kopargaon

India, kopargaon

Skills

- SQL
- Python
- ETL
- Machine Learning
- NoSQL
- Data Modeling
- Apache Spark
- Data Warehousing
- Big Data
- Data Mining
- Data Analysis
- Pyspark
- AWS
- Azure Data Factory
- Redshift
- Data Modeling
- Data Migration

Delivered curated Gold-layer datasets to Couchbase for ML consumption via Azure ML Studio

- Collaborated with ML teams to ensure data performance, usability, and readiness for cargo optimization models
- Supported scala

Data Engineer

Global Pharma /

Project Name

AU–EU Pharmaceutical Sales Data Pipeline Project

Technical Skills

Amazon S3, AWS Lambda, AWS Glue, Amazon Redshift, Amazon SNS, SFTP, Python, ETL, GDPR Compliance, Data Governance, Automation, CI/CD

Project Overview

A global pharmaceutical distributor faced challenges standardizing and analyzing sales data from **Australia (AU) and Europe (EU)** due to varied formats, multiple delivery channels, and region-specific compliance requirements (e.g., GDPR).

The initial unified pipeline created complexity, slowed processing, and caused single points of failure. The project re-architected the system into **dedicated AU and EU pipelines**, leveraging serverless AWS services and a region-aware ETL framework. Key features included automated ingestion, region-specific transformation, regulatory compliance, audit logging, and analytics-ready outputs. The solution enabled scalable, parallel execution of pipelines while maintaining traceability and operational efficiency.

Roles & Responsibilities

- Led redesign from a single unified pipeline to **region-specific AU and EU pipelines**
- Built **Lambda functions** for file ingestion, orchestration, and automated region detection
- Developed **AWS Glue ETL scripts** for AU-specific validations and EU GDPR-compliant transformations
- Designed **Amazon Redshift schemas** and ingestion processes for region-scoped analytics
- Implemented **audit logging, notifications (SNS), and monitoring** for operational transparency
- Coordinated with QA and DevOps teams for **CI/CD pipelines and automated deployments**
- Ensured data quality, regulatory compliance, and structured reporting for business stakeholders

Business Impact

- **90% reduction** in manual interventions and data corrections
- Parallel regional pipeline execution improved **reporting turnaround times**
- Strengthened **GDPR compliance** and audit traceability
- Delivered **clean, reliable datasets** for regional analytics and decision-making
- Established a **scalable framework** for onboarding additional regions (US, APAC)

Data Engineer

Kaiser Permanente /

Project Name

Kaiser Permanente – Healthcare Data Engineering Platform

Technical Skills

Azure Data Lake Storage (ADLS), Azure Databricks, Apache Spark, Spark SQL, Delta Lake, Azure Synapse, Azure Purview, Azure Key Vault, Azure Monitor, Azure Data Factory, Python, Pandas, DBT Cloud, Pytest, Infrastructure as Code (IaC), Medallion Architecture, Data Modeling, Schema Enforcement, Data Governance

Apache Hadoop	◆◆◆◆◆
Data Visualization	◆◆◆◆◆
Business Intelligence	◆◆◆◆◆
Cloud Computing	◆◆◆◆◆

Hobbies

- Bike riding
- Reading
- Traveling

Languages

- English
- Hindi
- Marathi

Project Overview

This project involved building a secure, scalable, and HIPAA-compliant healthcare data engineering platform for Kaiser Permanente. The organization was dealing with fragmented and unstructured healthcare data originating from patient–doctor interactions and medical records exposed through APIs. Manual data cleaning and inconsistent processing led to delayed insights, operational inefficiencies, and increased compliance risks.

The solution implemented a Medallion Architecture (Bronze–Silver–Gold) using Azure and Databricks to automate ingestion, validation, transformation, and delivery of healthcare data. The platform ensured strong data quality, lineage, governance, and auditability while enabling near real-time analytics and downstream machine learning use cases. Curated datasets improved accessibility for clinicians and laid the foundation for advanced analytics and AI-driven healthcare applications.

Roles

- Designed and implemented scalable, HIPAA-compliant data pipelines using Azure Databricks and Apache Spark
- Ingested raw healthcare JSON data from multiple APIs into ADLS with metadata capture for lineage and auditability
- Implemented Medallion Architecture (Bronze, Silver, Gold) with Delta Lake for ACID compliance and schema evolution
- Developed ETL workflows including JSON flattening, schema enforcement, deduplication, and null handling
- Built automated data quality and validation checks using DBT Cloud and Pytest (schema, null, and integrity checks)
- Orchestrated and monitored pipelines using Databricks Workflows and Azure Data Factory with retry and alert mechanisms

Business impact

- The project delivered a secure, compliant, and scalable healthcare data platform that transformed thats how Kaiser Permanente managed clinical data.
- Doctors gained access to accurate, curated patient information in near real time, reducing operational overhead and improving patient care
- The platform also laid the foundation for advanced analytics, predictive modeling, and LLM-powered clinical assistant use cases, while ensuring strong governance, security, and regulatory compliance